

In [ ]:

Predicting Custermer Churn for Bank Customers

Predicting Churn for Bank Customers

Libraries importing

```
In [1]: # Data Visualisation Libraries
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
sns.set_style('darkgrid')
```

```
In [2]: df = pd.read_csv(r"C:\Users\SSD\Downloads\Churn_Modelling.csv")
```

```
In [3]: df
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	B
0	1	15634602	Hargrave	619	France	Female	42	2	0.
1	2	15647311	Hill	608	Spain	Female	41	1	83
2	3	15619304	Onio	502	France	Female	42	8	15
3	4	15701354	Boni	699	France	Female	39	1	0.
4	5	15737888	Mitchell	850	Spain	Female	43	2	12
...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.
9996	9997	15569892	Johnstone	516	France	Male	35	10	57
9997	9998	15584532	Liu	709	France	Female	36	7	0.
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75
9999	10000	15628319	Walker	792	France	Female	28	4	13

10000 rows × 14 columns

```
In [4]: df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	Yes	Yes	100000.00	0
1	2	15647311	Hill	608	Spain	Female	41	1	83807.00	1	Yes	Yes	100000.00	0
2	3	15619304	Onio	502	France	Female	42	8	15966.00	1	Yes	Yes	100000.00	0
3	4	15701354	Boni	699	France	Female	39	1	0.00	1	Yes	Yes	100000.00	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	12557.00	1	Yes	Yes	100000.00	0

```
In [5]: df.shape
```

(10000, 14)

```
In [6]: df.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1, inplace=True)
```

```
df.columns
```

```
Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
       'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary',
       'Exited'],
      dtype='object')
```

```
In [7]: df.shape
```

(10000, 11)

In [8]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   CreditScore      10000 non-null   int64  
 1   Geography        10000 non-null   object  
 2   Gender            10000 non-null   object  
 3   Age               10000 non-null   int64  
 4   Tenure            10000 non-null   int64  
 5   Balance           10000 non-null   float64 
 6   NumOfProducts     10000 non-null   int64  
 7   HasCrCard         10000 non-null   int64  
 8   IsActiveMember    10000 non-null   int64  
 9   EstimatedSalary   10000 non-null   float64 
 10  Exited            10000 non-null   int64  
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB
```

In [9]: *#Thankfully, there are no missing values in our DataFrame.*

The describe() method gives us a statistical summary of the numerical features:

In [10]: df.describe().T

	count	mean	std	min	25%	50%	75%	max
CreditScore	10000.0	650.528800	96.653299	350.00	584.00	652.000	718.0000	85
Age	10000.0	38.921800	10.487806	18.00	32.00	37.000	44.0000	92
Tenure	10000.0	5.012800	2.892174	0.00	3.00	5.000	7.0000	10
Balance	10000.0	76485.889288	62397.405202	0.00	0.00	97198.540	127644.2400	25
NumOfProducts	10000.0	1.530200	0.581654	1.00	1.00	1.000	2.0000	4.0
HasCrCard	10000.0	0.705500	0.455840	0.00	0.00	1.000	1.0000	1.0
IsActiveMember	10000.0	0.515100	0.499797	0.00	0.00	1.000	1.0000	1.0
EstimatedSalary	10000.0	100090.239881	57510.492818	11.58	51002.11	100193.915	149388.2475	19
Exited	10000.0	0.203700	0.402769	0.00	0.00	0.000	0.0000	1.0

In [11]: *#Cheking for missing numbers*

```
In [12]: df.isnull().sum()
```

```
CreditScore      0
Geography       0
Gender          0
Age             0
Tenure          0
Balance          0
NumOfProducts    0
HasCrCard        0
IsActiveMember   0
EstimatedSalary  0
Exited           0
dtype: int64
```

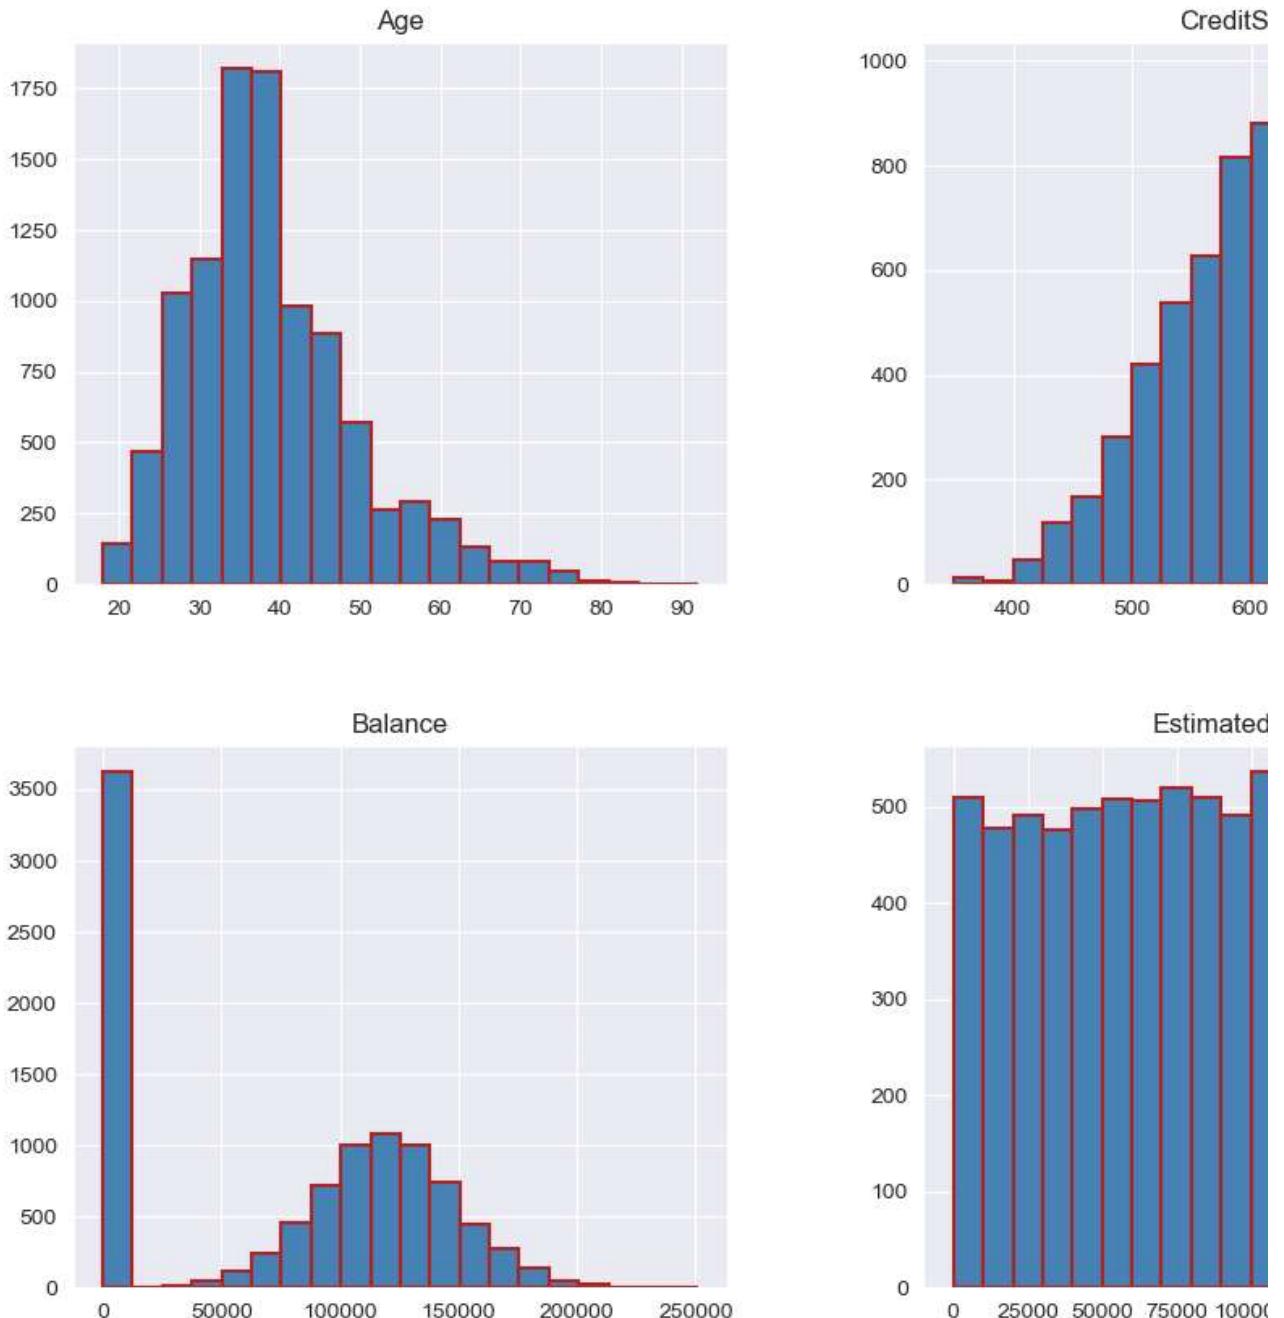
## Visulization

```
In [13]: continuous = ['Age', 'CreditScore', 'Balance', 'EstimatedSalary']
categorical = ['Geography', 'Gender', 'Tenure', 'NumOfProducts', 'HasCrCard', '']

print('Continuous: ', ', '.join(continuous))
print('Categorical: ', ', '.join(categorical))

Continuous: Age, CreditScore, Balance, EstimatedSalary
Categorical: Geography, Gender, Tenure, NumOfProducts, HasCrCard, IsActiveMember
```

```
In [14]: df[continuous].hist(figsize=(12, 10),
                         bins=20,
                         layout=(2, 2),
                         color='steelblue',
                         edgecolor='firebrick',
                         linewidth=1.5);
```



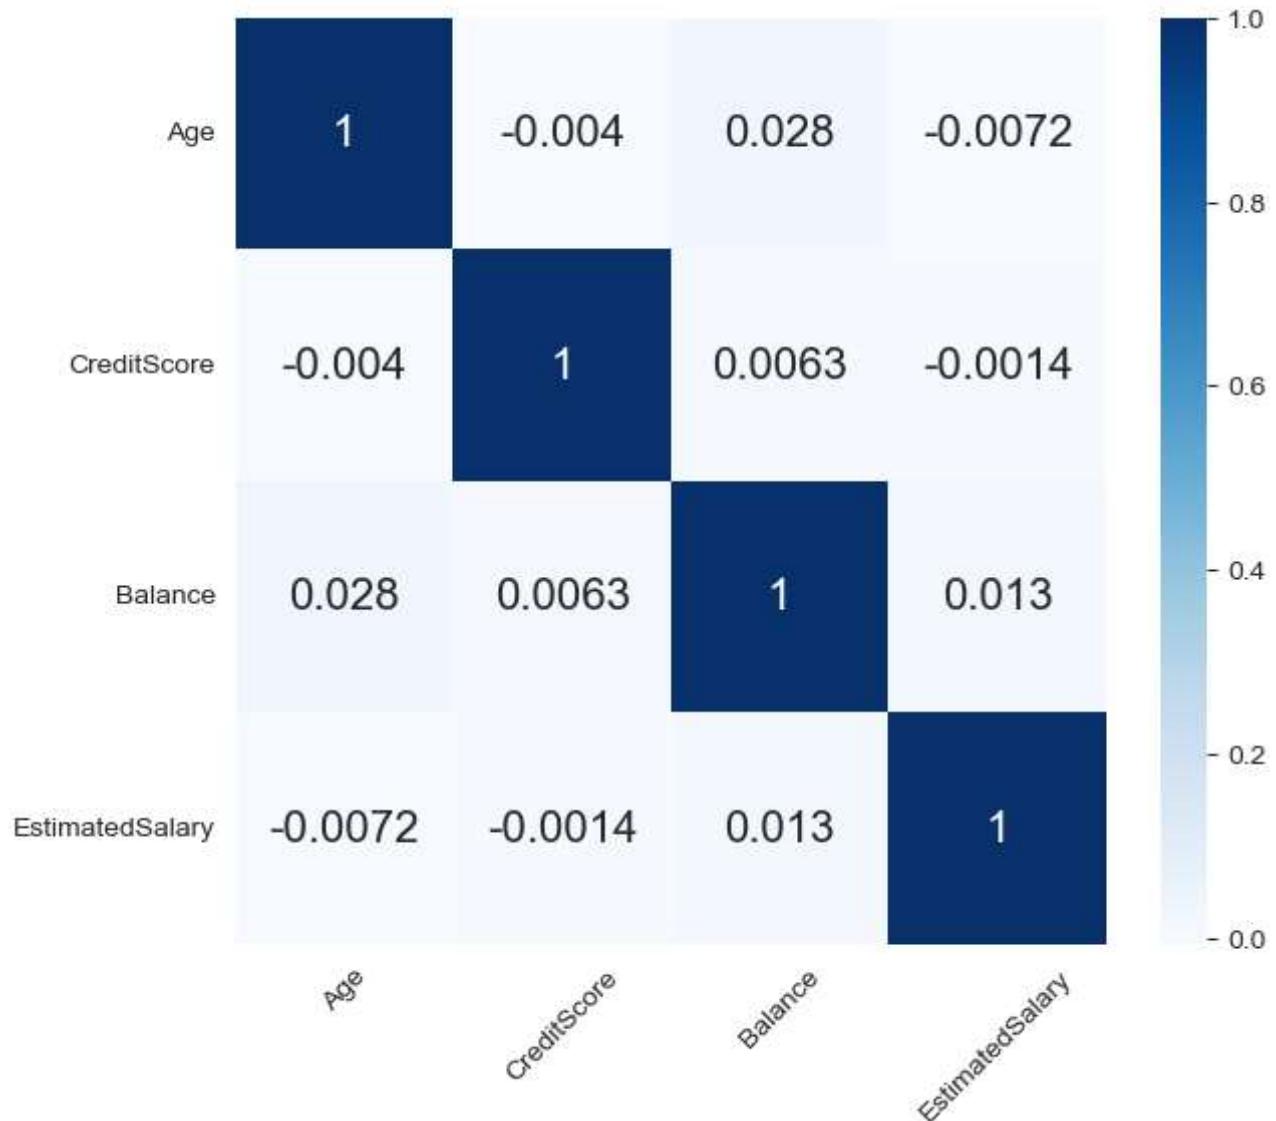
'Age' is slightly tail-heavy, i.e. it extends more further to the right of the median than 'CreditScore' are above 600, If we ignore the first bin, 'Balance' follows a fairly normal distribution of 'EstimatedSalary' is more or less uniform and provides little information.

## Looking for Correlations

```
In [15]: fig, ax = plt.subplots(figsize=(7, 6))

sns.heatmap(df[continuous].corr(),
            annot=True,
            annot_kws={'fontsize': 16},
            cmap='Blues',
            ax=ax)

ax.tick_params(axis='x', rotation=45)
ax.tick_params(axis='y', rotation=360);
```



```
In [16]: # Create subsets
```

```
df_churned = df[df['Exited'] == 1]
df_retained = df[df['Exited'] == 0]
```

```
In [17]: df_churned
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard
0	619	France	Female	42	2	0.00	1	1
2	502	France	Female	42	8	159660.80	3	1
5	645	Spain	Male	44	8	113755.78	2	1
7	376	Germany	Female	29	4	115046.74	4	1
16	653	Germany	Male	58	1	132602.88	1	1
...	...	...	...	...	...	...	...	...
9981	498	Germany	Male	42	3	152039.70	1	1
9982	655	Germany	Female	46	7	137145.12	1	1
9991	597	France	Female	53	4	88381.21	1	1
9997	709	France	Female	36	7	0.00	1	0
9998	772	Germany	Male	42	3	75075.31	2	1

2037 rows × 11 columns

```
In [18]: df_retained
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard
1	608	Spain	Female	41	1	83807.86	1	0
3	699	France	Female	39	1	0.00	2	0
4	850	Spain	Female	43	2	125510.82	1	1
6	822	France	Male	50	7	0.00	2	1
8	501	France	Male	44	4	142051.07	2	0
...	...	...	...	...	...	...	...	...
9993	644	France	Male	28	7	155060.41	1	1
9994	800	France	Female	29	2	0.00	2	0
9995	771	France	Male	39	5	0.00	2	1
9996	516	France	Male	35	10	57369.61	1	1
9999	792	France	Female	28	4	130142.79	1	1

7963 rows × 11 columns

```
In [19]: # Plotting

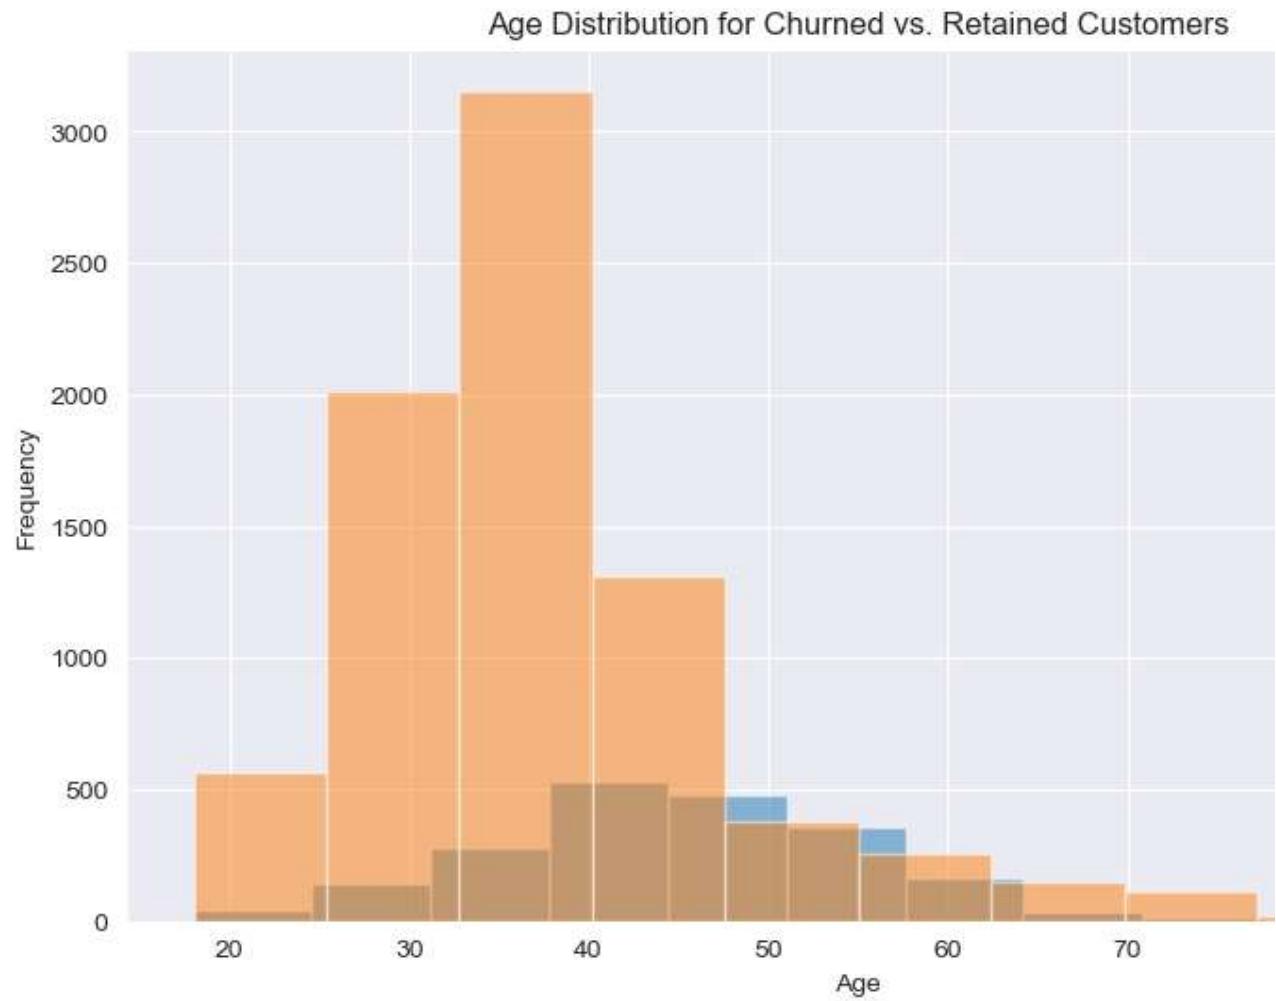
plt.figure(figsize=(10, 6))

# Plot age distribution for churned customers
plt.hist(df_churned['Age'], bins=10, alpha=0.5, label='Churned')

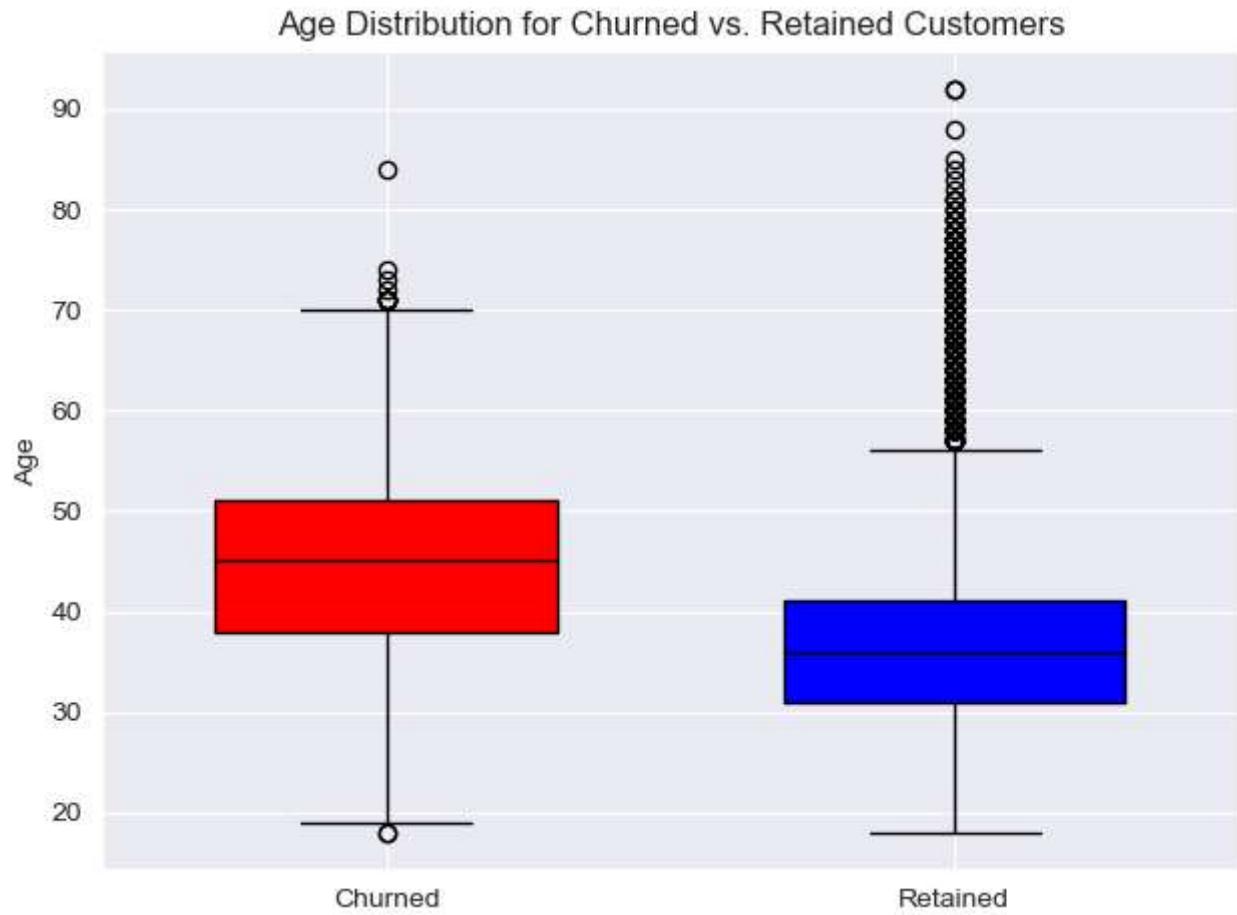
# Plot age distribution for retained customers
plt.hist(df_retained['Age'], bins=10, alpha=0.5, label='Retained')

plt.title('Age Distribution for Churned vs. Retained Customers')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.legend()
plt.grid(True)

# Show the plot
plt.show()
```



```
In [20]: # Create boxplot for churned customers' credit scores  
plt.boxplot(df_churned['Age'], positions=[1], widths=0.6, patch_artist=True, boxprops=dict(facecolor='red'))  
  
# Create boxplot for retained customers' credit scores  
plt.boxplot(df_retained['Age'], positions=[2], widths=0.6, patch_artist=True, boxprops=dict(facecolor='blue'))  
  
# Adjust x-axis labels  
plt.xticks([1, 2], ['Churned', 'Retained'])  
  
plt.title('Age Distribution for Churned vs. Retained Customers')  
plt.ylabel('Age')  
plt.grid(True)  
  
# Show the plot  
plt.tight_layout()  
plt.show()
```



```
In [21]: # Plotting

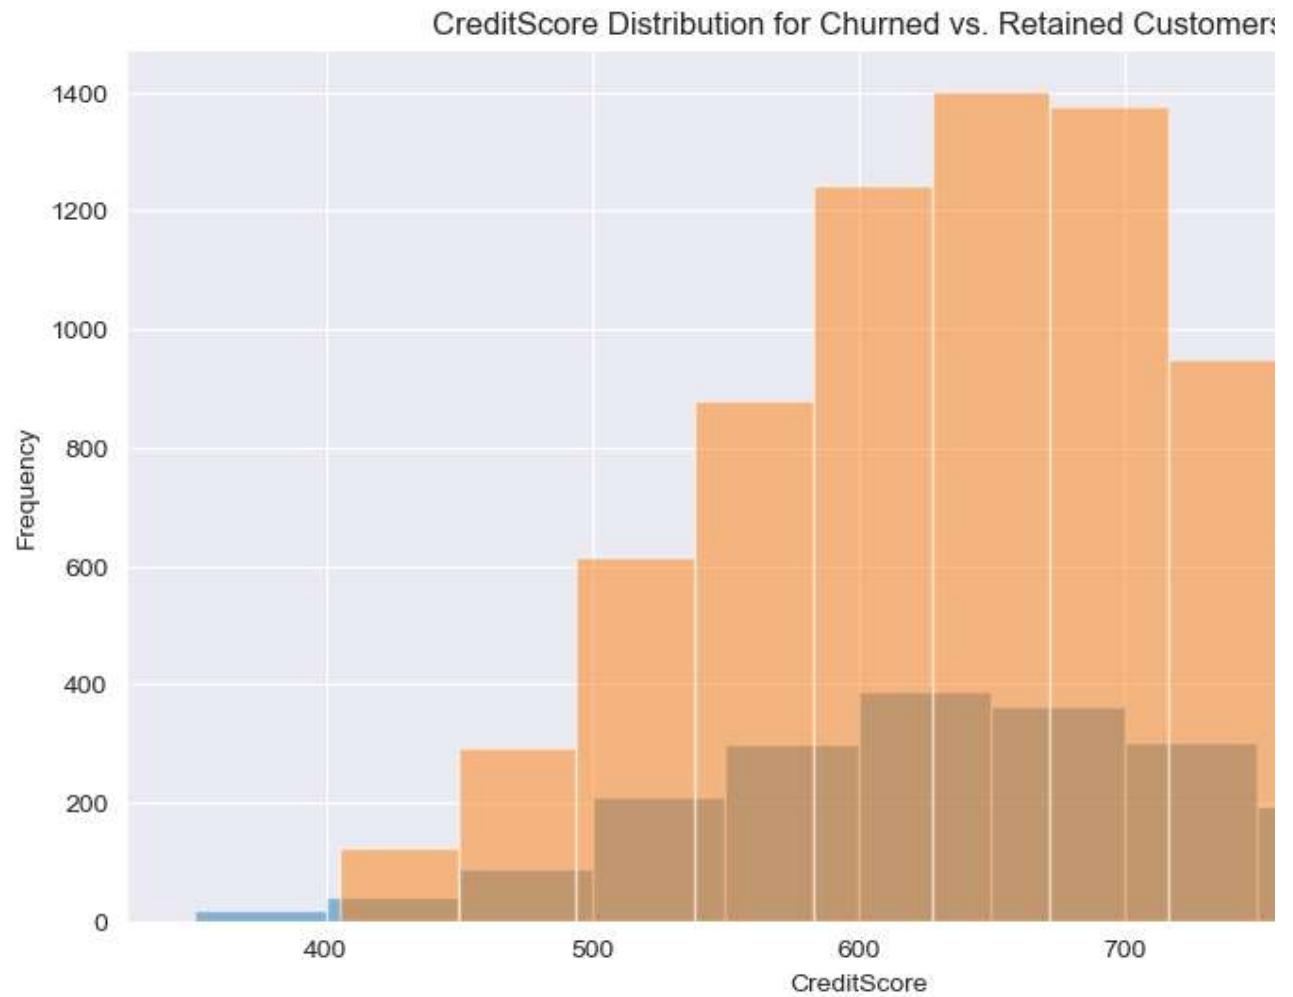
plt.figure(figsize=(10, 6))

# Plot age distribution for churned customers
plt.hist(df_churned['CreditScore'], bins=10, alpha=0.5, label='Churned')

# Plot age distribution for retained customers
plt.hist(df_retained['CreditScore'], bins=10, alpha=0.5, label='Retained')

plt.title('CreditScore Distribution for Churned vs. Retained Customers')
plt.xlabel('CreditScore')
plt.ylabel('Frequency')
plt.legend()
plt.grid(True)

# Show the plot
plt.show()
```



```
In [22]: # Create boxplot for churned customers' credit scores  
plt.boxplot(df_churned['CreditScore'], positions=[1], widths=0.6, patch_artist=1  
  
# Create boxplot for retained customers' credit scores  
plt.boxplot(df_retained['CreditScore'], positions=[2], widths=0.6, patch_artist=1  
  
# Adjust x-axis labels  
plt.xticks([1, 2], ['Churned', 'Retained'])  
  
plt.title('CreditScore Distribution for Churned vs. Retained Customers')  
plt.ylabel('CreditScore')  
plt.grid(True)  
  
# Show the plot  
plt.tight_layout()  
plt.show()
```



```
In [23]: # Plotting

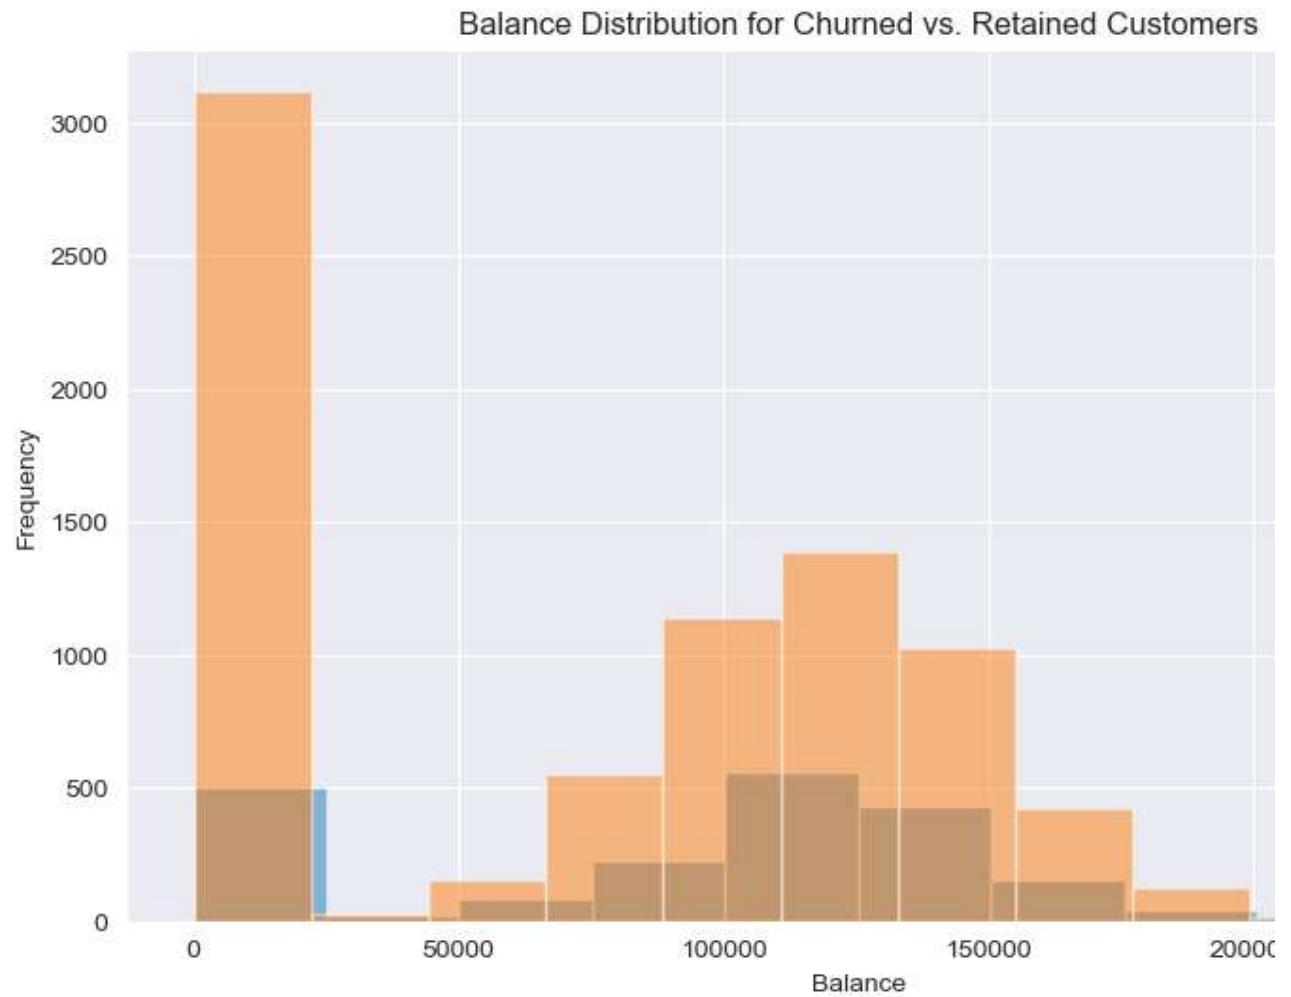
plt.figure(figsize=(10, 6))

# Plot age distribution for churned customers
plt.hist(df_churned['Balance'], bins=10, alpha=0.5, label='Churned')

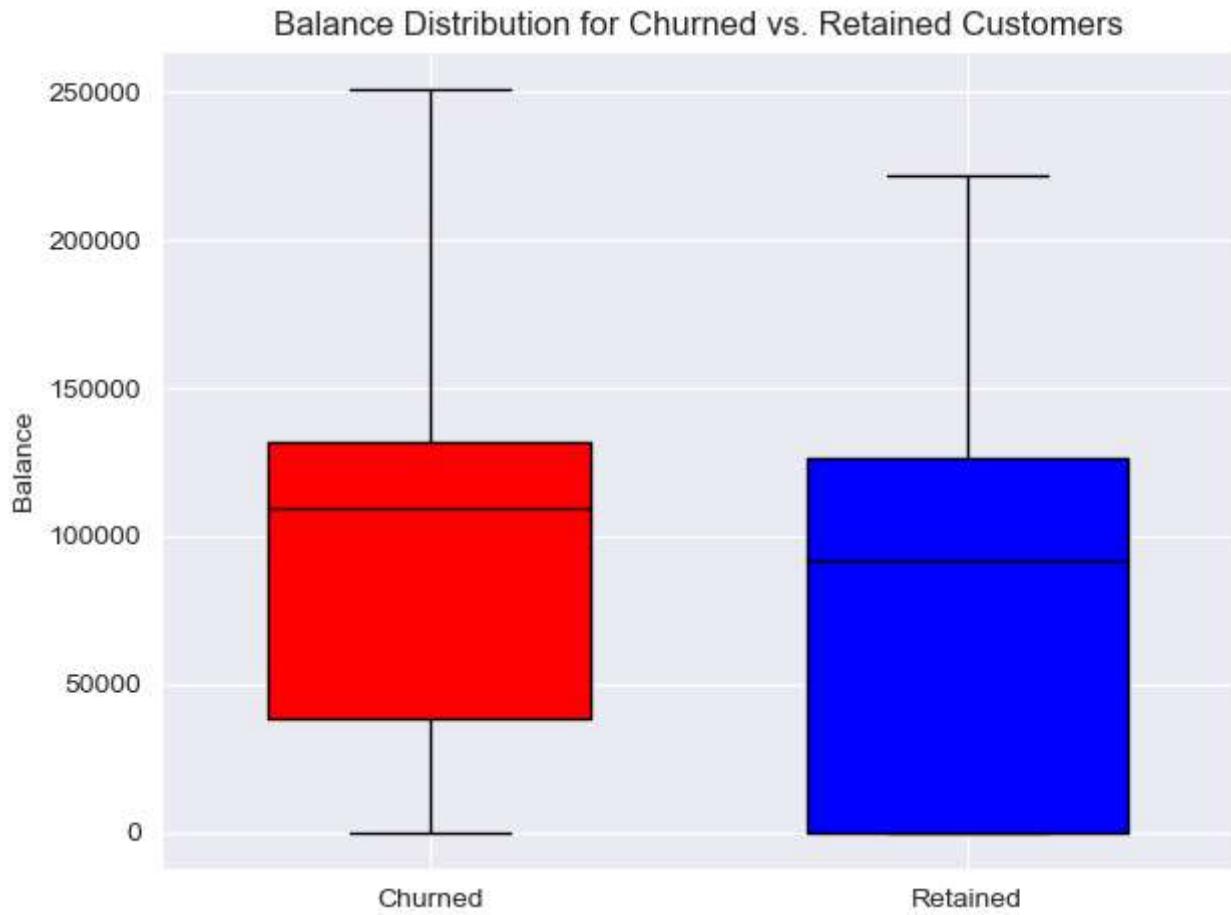
# Plot age distribution for retained customers
plt.hist(df_retained['Balance'], bins=10, alpha=0.5, label='Retained')

plt.title('Balance Distribution for Churned vs. Retained Customers')
plt.xlabel('Balance')
plt.ylabel('Frequency')
plt.legend()
plt.grid(True)

# Show the plot
plt.show()
```



```
In [24]: # Create boxplot for churned customers' credit scores  
plt.boxplot(df_churned['Balance'], positions=[1], widths=0.6, patch_artist=True,  
  
# Create boxplot for retained customers' credit scores  
plt.boxplot(df_retained['Balance'], positions=[2], widths=0.6, patch_artist=True)  
  
# Adjust x-axis labels  
plt.xticks([1, 2], ['Churned', 'Retained'])  
  
plt.title('Balance Distribution for Churned vs. Retained Customers')  
plt.ylabel('Balance')  
plt.grid(True)  
  
# Show the plot  
plt.tight_layout()  
plt.show()
```



```
In [25]: # Plotting

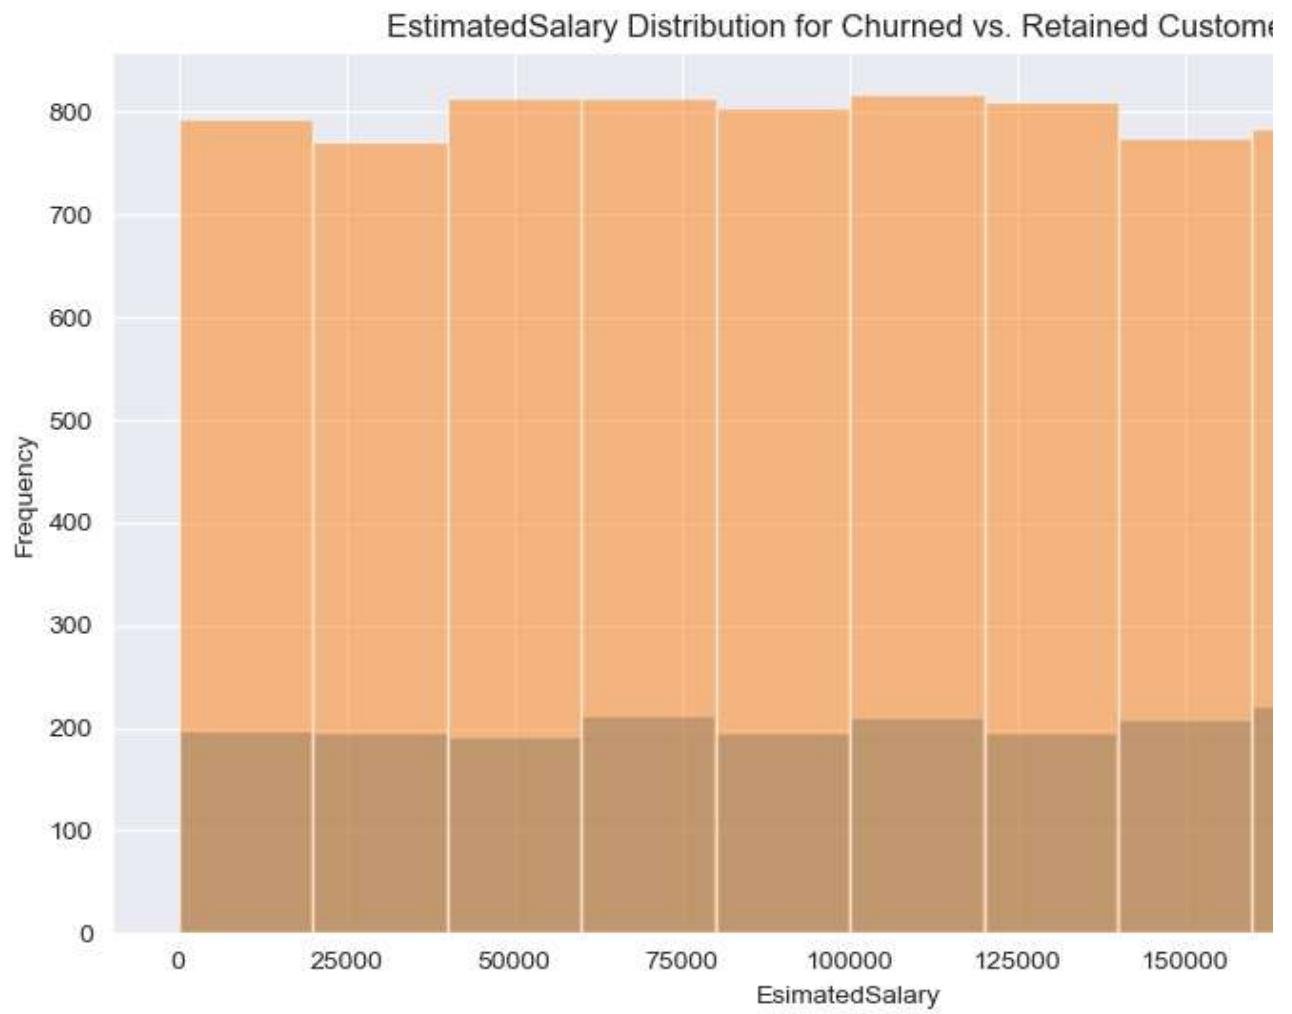
plt.figure(figsize=(10, 6))

# Plot age distribution for churned customers
plt.hist(df_churned['EstimatedSalary'], bins=10, alpha=0.5, label='Churned')

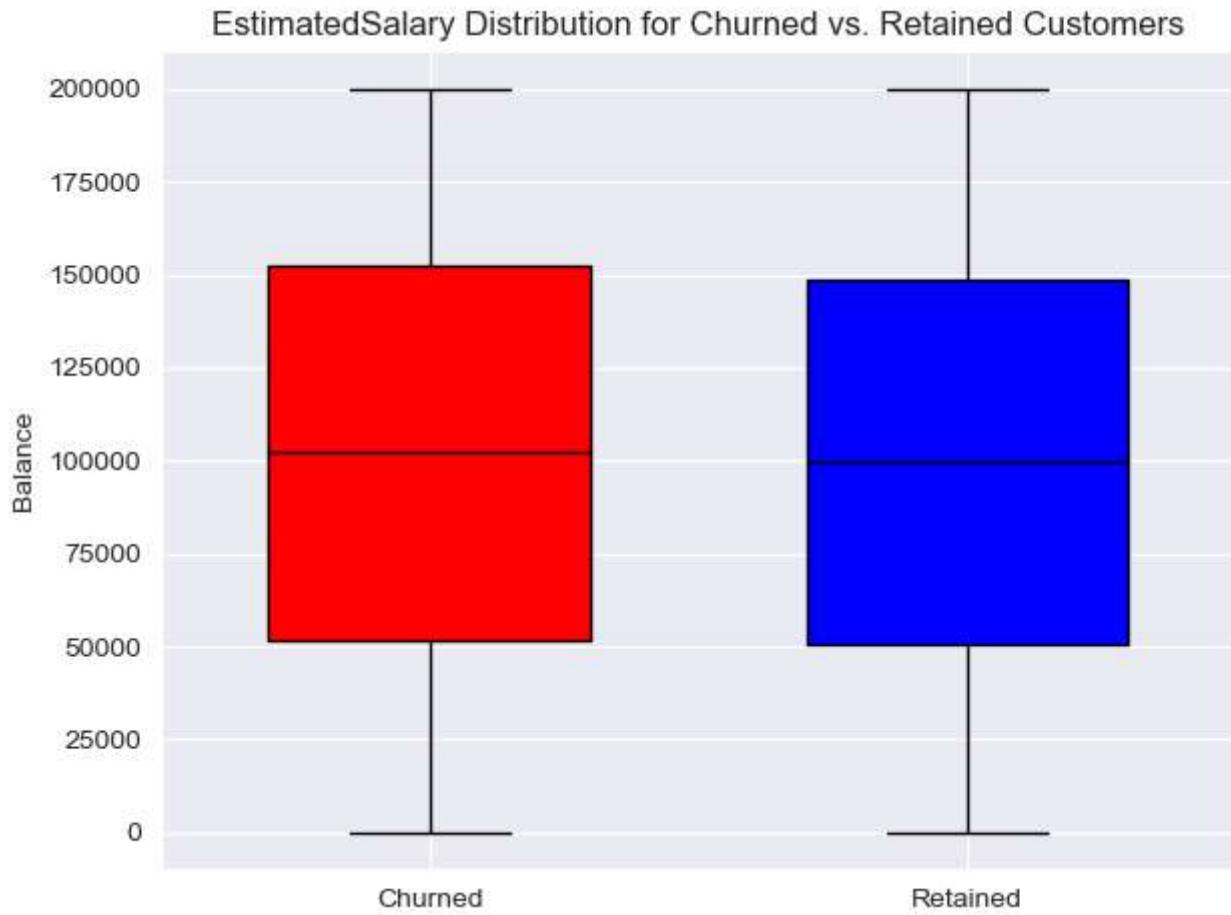
# Plot age distribution for retained customers
plt.hist(df_retained['EstimatedSalary'], bins=10, alpha=0.5, label='Retained')

plt.title('EstimatedSalary Distribution for Churned vs. Retained Customers')
plt.xlabel('EstimatedSalary')
plt.ylabel('Frequency')
plt.legend()
plt.grid(True)

# Show the plot
plt.show()
```



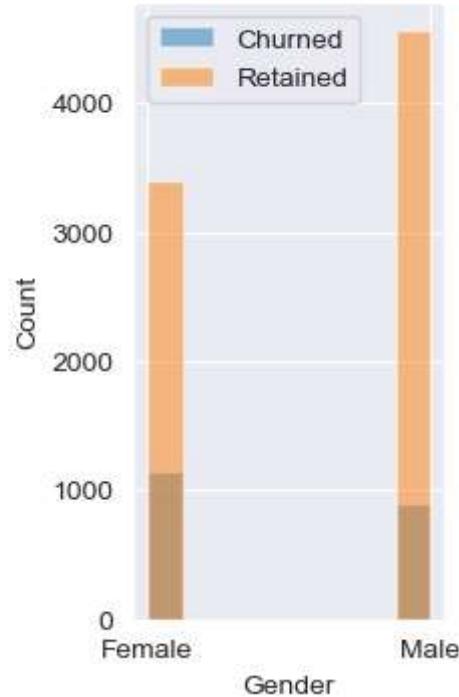
```
In [26]: # Create boxplot for churned customers' credit scores  
plt.boxplot(df_churned['EstimatedSalary'], positions=[1], widths=0.6, patch_arti  
  
# Create boxplot for retained customers' credit scores  
plt.boxplot(df_retained['EstimatedSalary'], positions=[2], widths=0.6, patch_arti  
  
# Adjust x-axis labels  
plt.xticks([1, 2], ['Churned', 'Retained'])  
  
plt.title('EstimatedSalary Distribution for Churned vs. Retained Customers')  
plt.ylabel('Balance')  
plt.grid(True)  
  
# Show the plot  
plt.tight_layout()  
plt.show()
```



## Categorical Variables

```
In [27]: # Plotting  
plt.figure(figsize=(2, 4))  
  
# Plot age distribution for churned customers  
plt.hist(df_churned['Gender'], bins=8, alpha=0.5, label='Churned')  
  
# Plot age distribution for retained customers  
plt.hist(df_retained['Gender'], bins=8, alpha=0.5, label='Retained')  
  
plt.title('Gender Distribution for Churned vs. Retained Customers')  
plt.xlabel('Gender')  
plt.ylabel('Count')  
plt.legend()  
plt.grid(True)  
  
# Show the plot  
plt.show()
```

Gender Distribution for Churned vs. Retained Customers



```
In [28]: import warnings  
warnings.simplefilter(action='ignore', category=FutureWarning)
```

---

```
In [29]: df_cat = df[categorical]

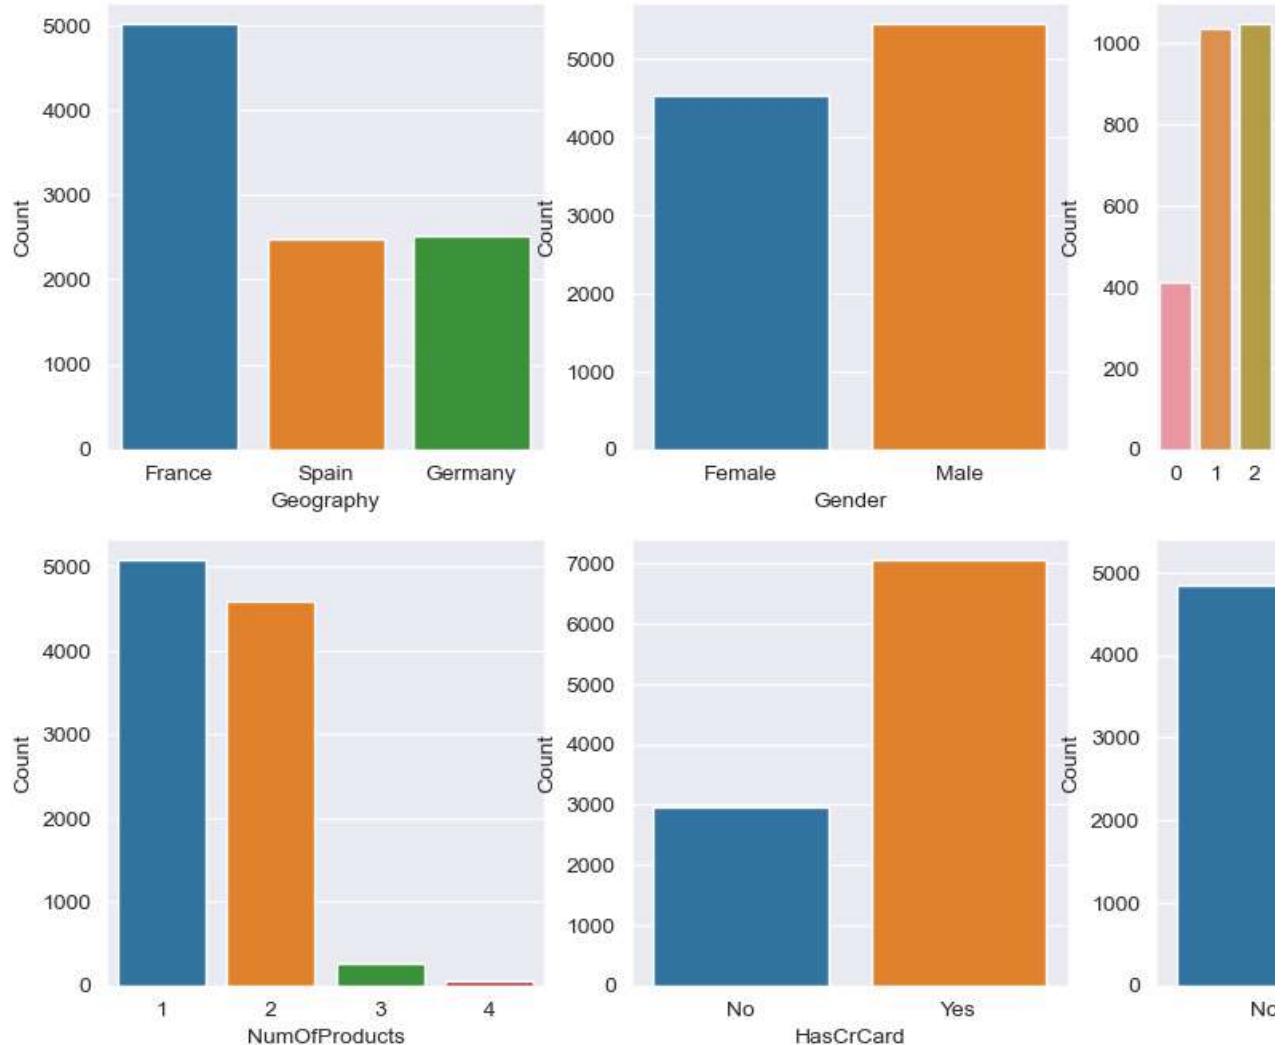
fig, ax = plt.subplots(2, 3, figsize=(12, 8))

for index, column in enumerate(df_cat.columns):

    plt.subplot(2, 3, index + 1)
    sns.countplot(x=column, data=df)

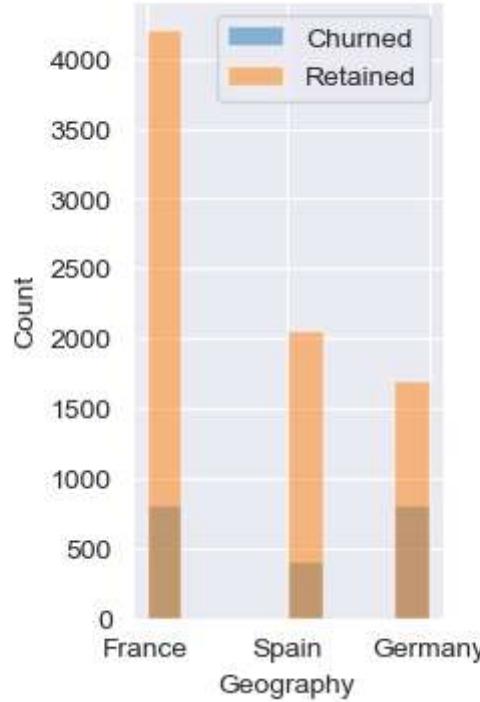
    plt.ylabel('Count')
    if (column == 'HasCrCard' or column == 'IsActiveMember'):
        plt.xticks([0, 1], ['No', 'Yes'])


```



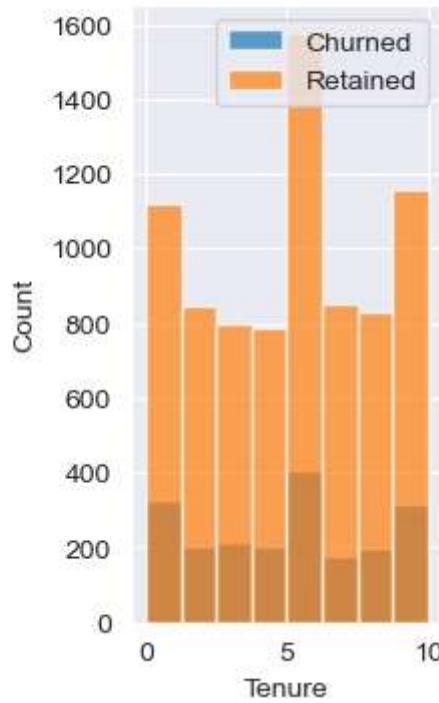
```
In [30]: # Plotting  
plt.figure(figsize=(2, 4))  
  
# Plot age distribution for churned customers  
plt.hist(df_churned['Geography'], bins=8, alpha=0.5, label='Churned')  
  
# Plot age distribution for retained customers  
plt.hist(df_retained['Geography'], bins=8, alpha=0.5, label='Retained')  
  
plt.title('Geography Distribution for Churned vs. Retained Customers')  
plt.xlabel('Geography')  
plt.ylabel('Count')  
plt.legend()  
plt.grid(True)  
  
# Show the plot  
plt.show()
```

Geography Distribution for Churned vs. Retained Customers



```
In [31]: # Plotting  
plt.figure(figsize=(2, 4))  
  
# Plot age distribution for churned customers  
plt.hist(df_churned['Tenure'], bins=8, alpha=0.7, label='Churned')  
  
# Plot age distribution for retained customers  
plt.hist(df_retained['Tenure'], bins=8, alpha=0.7, label='Retained')  
  
plt.title('Tenure Distribution for Churned vs. Retained Customers')  
plt.xlabel('Tenure')  
plt.ylabel('Count')  
plt.legend()  
plt.grid(True)  
  
# Show the plot  
plt.show()
```

Tenure Distribution for Churned vs. Retained Customers



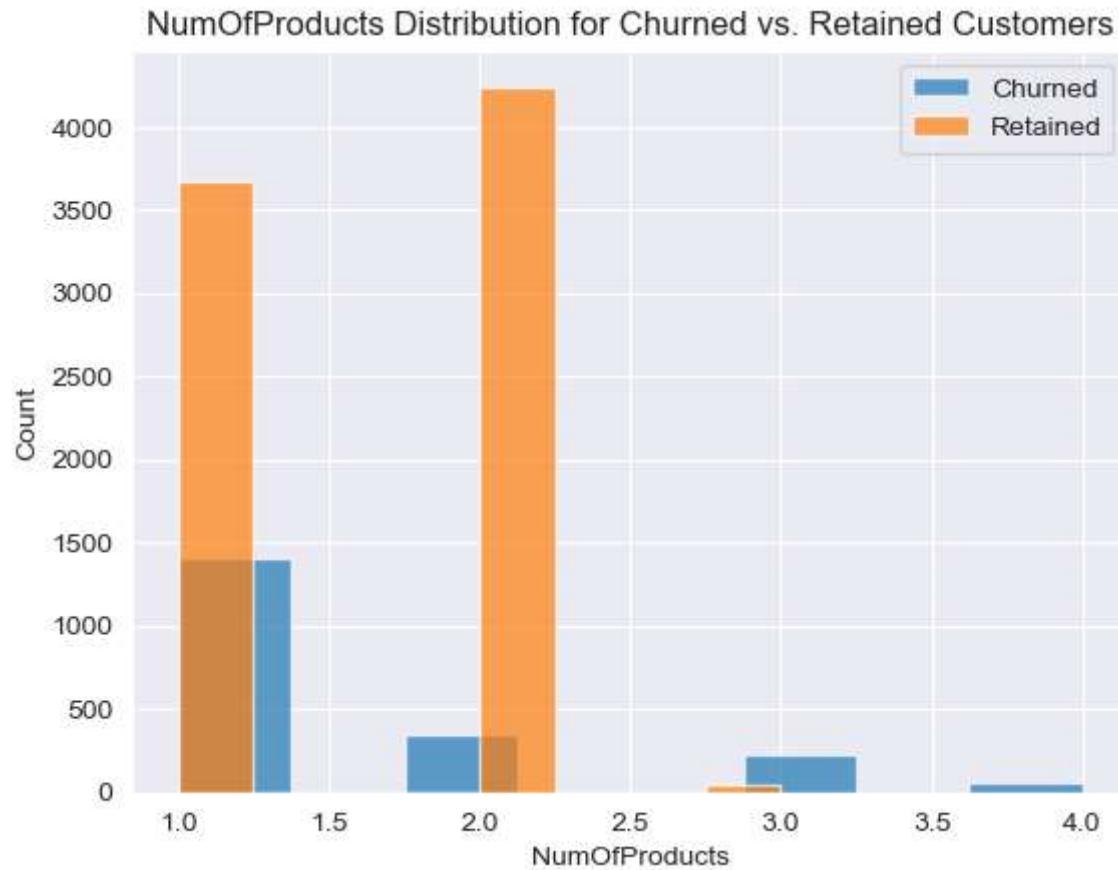
```
In [32]: # Plotting
plt.figure()

# Plot age distribution for churned customers
plt.hist(df_churned[ 'NumOfProducts' ], bins=8, alpha=0.7, label='Churned')

# Plot age distribution for retained customers
plt.hist(df_retained[ 'NumOfProducts' ], bins=8, alpha=0.7, label='Retained')

plt.title('NumOfProducts Distribution for Churned vs. Retained Customers')
plt.xlabel('NumOfProducts')
plt.ylabel('Count')
plt.legend()
plt.grid(True)

# Show the plot
plt.show()
```

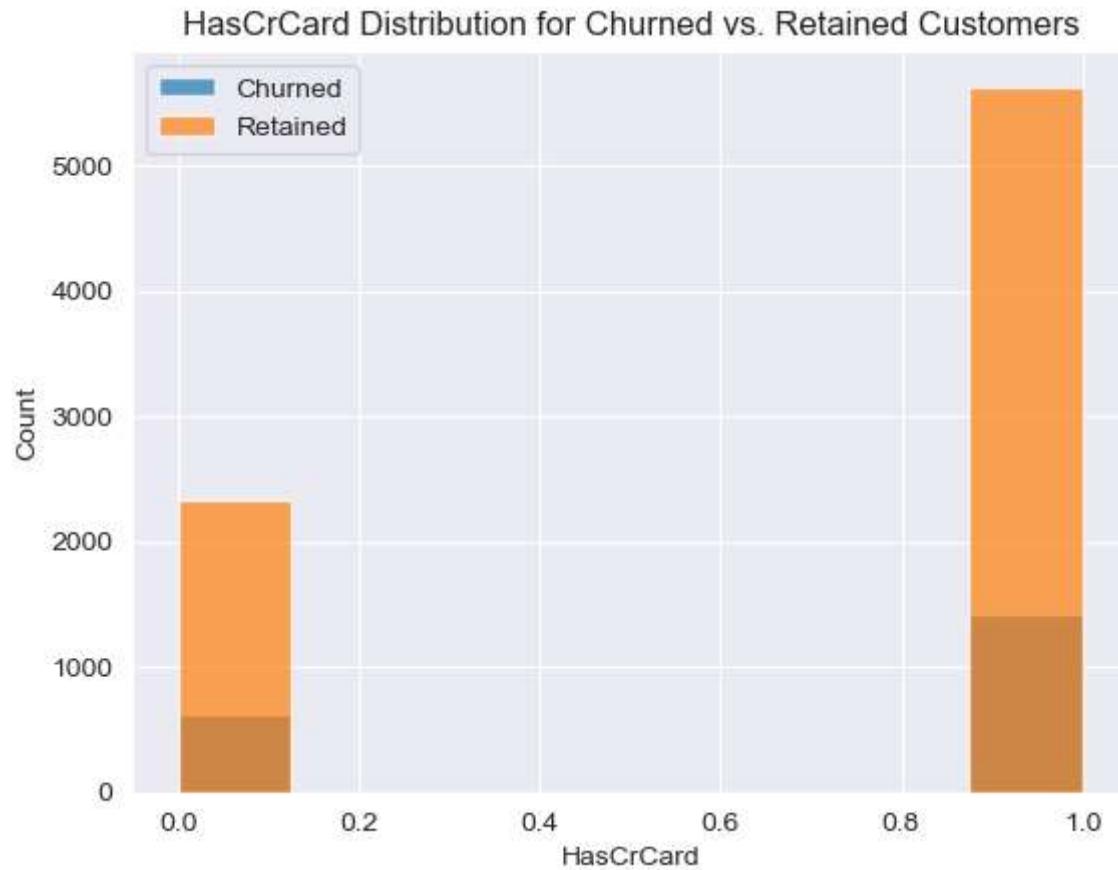


```
In [33]: # Plotting
plt.figure()

# Plot age distribution for churned customers
plt.hist(df_churned['HasCrCard'], bins=8, alpha=0.7, label='Churned')

# Plot age distribution for retained customers
plt.hist(df_retained['HasCrCard'], bins=8, alpha=0.7, label='Retained')

plt.title('HasCrCard Distribution for Churned vs. Retained Customers')
plt.xlabel('HasCrCard')
plt.ylabel('Count')
plt.legend()
plt.grid(True)
```



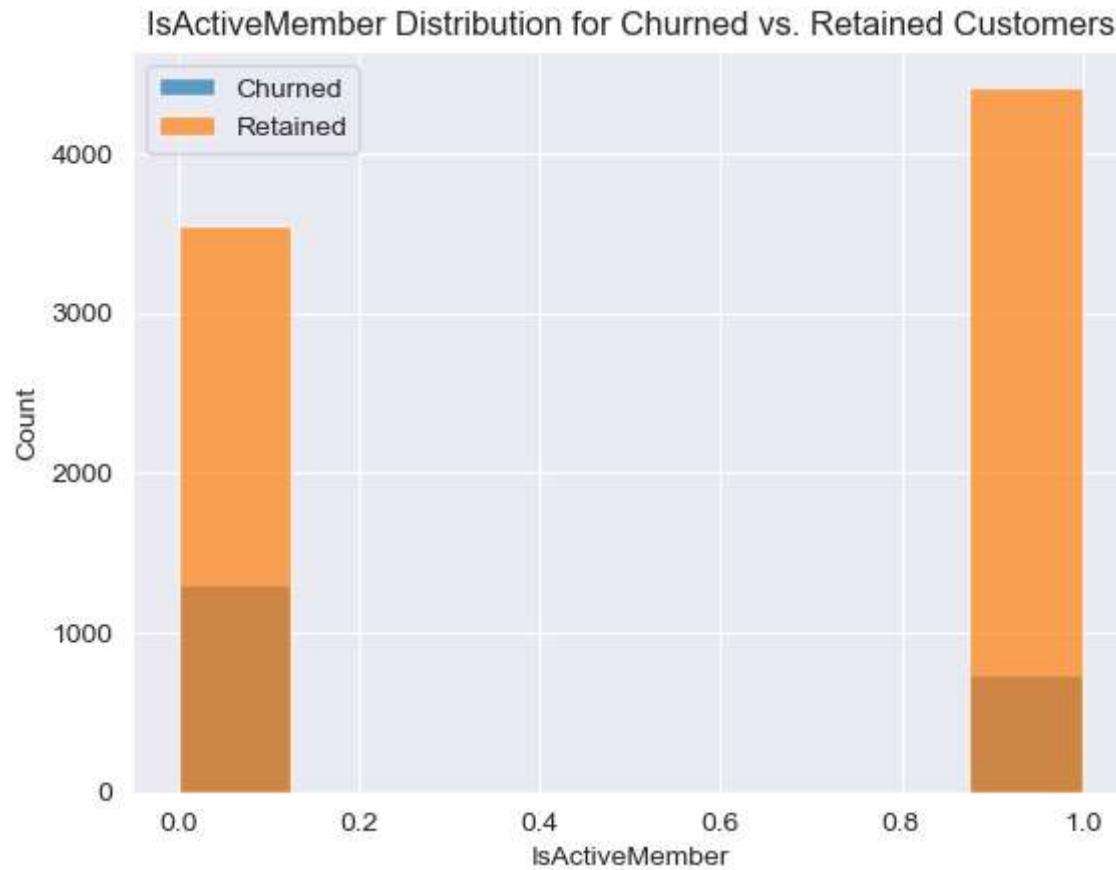
```
In [34]: # Plotting
plt.figure()

# Plot age distribution for churned customers
plt.hist(df_churned['IsActiveMember'], bins=8, alpha=0.7, label='Churned')

# Plot age distribution for retained customers
plt.hist(df_retained['IsActiveMember'], bins=8, alpha=0.7, label='Retained')

plt.title('IsActiveMember Distribution for Churned vs. Retained Customers')
plt.xlabel('IsActiveMember')
plt.ylabel('Count')
plt.legend()
plt.grid(True)

# Show the plot
plt.show()
```



```
In [35]: df.drop(['Tenure', 'HasCrCard', 'EstimatedSalary'], axis=1, inplace=True)  
df.columns
```

```
Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Balance', 'NumOfProducts',  
       'IsActiveMember', 'Exited'],  
      dtype='object')
```

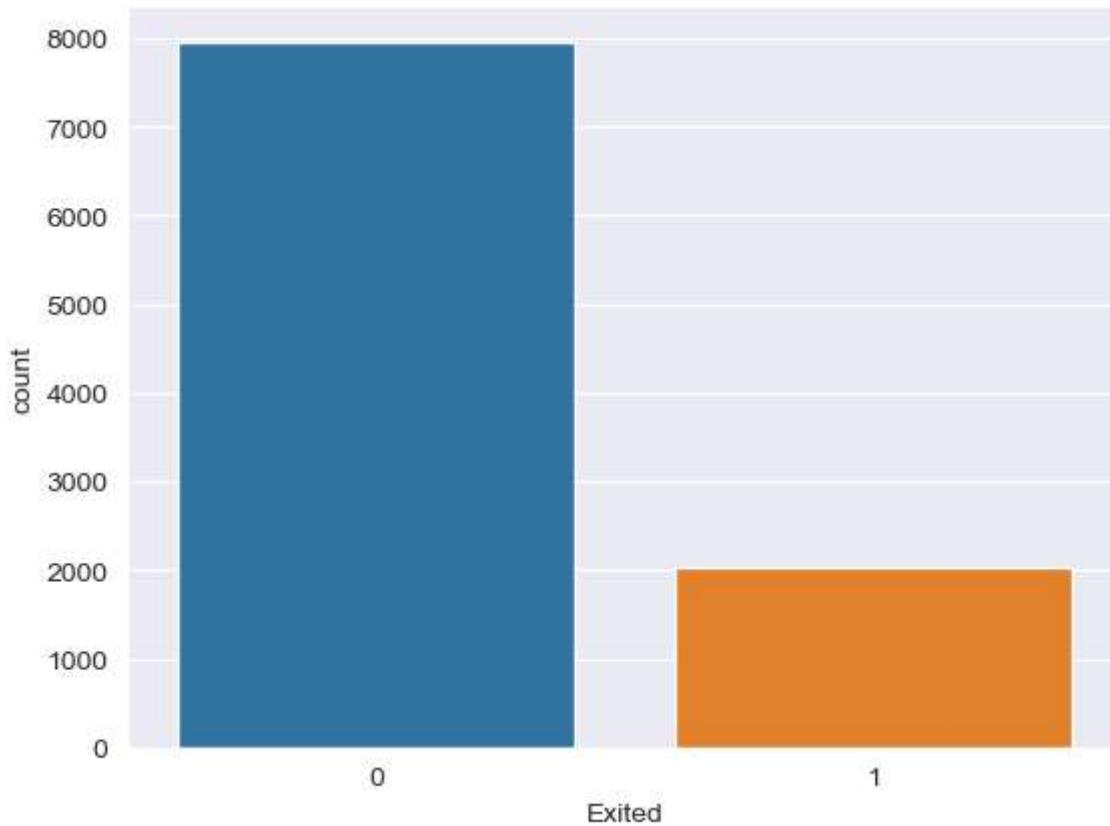
```
In [36]: df.shape
```

```
(10000, 8)
```

```
In [37]: df['Exited'].value_counts()
```

```
0    7963  
1    2037  
Name: Exited, dtype: int64
```

```
In [38]: # Assuming df is your DataFrame  
sns.countplot(x='Exited', data=df)  
  
# Optionally, show the plot  
plt.show()
```



split the featurevariable and tergetvariable

```
In [39]: X = df.iloc[:, 0:8].values  
y = df.iloc[:, -1].values
```

```
In [40]: X  
  
array([[619, 'France', 'Female', ..., 1, 1, 1],  
       [608, 'Spain', 'Female', ..., 1, 1, 0],  
       [502, 'France', 'Female', ..., 3, 0, 1],  
       ...,  
       [709, 'France', 'Female', ..., 1, 1, 1],  
       [772, 'Germany', 'Male', ..., 2, 0, 1],  
       [792, 'France', 'Female', ..., 1, 0, 0]], dtype=object)
```

```
In [41]: X.shape
```

```
(10000, 8)
```

```
In [42]: y
```

```
array([1, 0, 1, ..., 1, 1, 0], dtype=int64)
```

```
In [43]: y.shape
```

```
(10000,)
```

## Encoding categorical data

```
In [44]: from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()  
X[:, 2] = le.fit_transform(X[:, 2])
```

```
In [45]: X
```

```
array([[619, 'France', 0, ..., 1, 1, 1],  
       [608, 'Spain', 0, ..., 1, 1, 0],  
       [502, 'France', 0, ..., 3, 0, 1],  
       ...,  
       [709, 'France', 0, ..., 1, 1, 1],  
       [772, 'Germany', 1, ..., 2, 0, 1],  
       [792, 'France', 0, ..., 1, 0, 0]], dtype=object)
```

```
In [46]: from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()  
X[:, 1] = le.fit_transform(X[:, 1])
```

```
In [47]: X
```

```
array([[619, 0, 0, ..., 1, 1, 1],  
       [608, 2, 0, ..., 1, 1, 0],  
       [502, 0, 0, ..., 3, 0, 1],  
       ...,  
       [709, 0, 0, ..., 1, 1, 1],  
       [772, 1, 1, ..., 2, 0, 1],  
       [792, 0, 0, ..., 1, 0, 0]], dtype=object)
```

```
In [48]: y
```

```
array([1, 0, 1, ..., 1, 1, 0], dtype=int64)
```

```
In [49]: # Splitting the dataset into the Training set and Test set
from sklearn.model_selection import train_test_split
#for this observation let me selected as 100 observaion for test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20,random_state=42)
```

```
In [50]: X_train.shape
```

```
(8000, 8)
```

```
In [51]: X_test.shape
```

```
(2000, 8)
```

```
In [52]: y_train.shape
```

```
(8000,)
```

```
In [53]: y_test.shape
```

```
(2000,)
```

```
In [54]: X_train
```

```
array([[667, 2, 0, ..., 2, 0, 0],
       [427, 1, 1, ..., 1, 1, 0],
       [535, 0, 0, ..., 1, 0, 0],
       ...,
       [738, 0, 1, ..., 2, 0, 0],
       [590, 2, 0, ..., 2, 1, 0],
       [623, 1, 0, ..., 1, 0, 1]], dtype=object)
```

```
In [55]: from sklearn.preprocessing import StandardScaler  
sc = StandardScaler()  
  
X_train = sc.fit_transform(X_train)  
X_test = sc.transform(X_test)  
X_train  
  
array([[ 0.16958176,  1.51919821, -1.09168714, ...,  0.8095029 ,  
       -1.03227043, -0.50624244],  
      [-2.30455945,  0.3131264 ,  0.91601335, ..., -0.92159124,  
       0.9687384 , -0.50624244],  
      [-1.19119591, -0.89294542, -1.09168714, ..., -0.92159124,  
       -1.03227043, -0.50624244],  
      ...,  
      [ 0.9015152 , -0.89294542,  0.91601335, ...,  0.8095029 ,  
       -1.03227043, -0.50624244],  
      [-0.62420521,  1.51919821, -1.09168714, ...,  0.8095029 ,  
       0.9687384 , -0.50624244],  
      [-0.28401079,  0.3131264 , -1.09168714, ..., -0.92159124,  
       -1.03227043,  1.97533814]])
```

```
In [56]: X_test  
  
array([[-0.55204276,  0.3131264 , -1.09168714, ..., -0.92159124,  
       0.9687384 , -0.50624244],  
      [-1.31490297, -0.89294542, -1.09168714, ..., -0.92159124,  
       -1.03227043,  1.97533814],  
      [ 0.57162971,  1.51919821, -1.09168714, ..., -0.92159124,  
       0.9687384 , -0.50624244],  
      ...,  
      [-0.74791227,  1.51919821,  0.91601335, ...,  0.8095029 ,  
       -1.03227043, -0.50624244],  
      [-0.00566991,  0.3131264 ,  0.91601335, ..., -0.92159124,  
       0.9687384 , -0.50624244],  
      [-0.79945688,  0.3131264 ,  0.91601335, ..., -0.92159124,  
       -1.03227043, -0.50624244]])
```

Training XGBoost on the Training set

```
In [57]: from xgboost import XGBClassifier  
classifier = XGBClassifier()  
classifier.fit(X_train, y_train)  
  
XGBClassifier(base_score=None, booster=None, callbacks=None,  
              colsample_bylevel=None, colsample_bynode=None,  
              colsample_bytree=None, device=None, early_stopping_rounds=None,  
              enable_categorical=False, eval_metric=None, feature_types=None,  
              gamma=None, grow_policy=None, importance_type=None,  
              interaction_constraints=None, learning_rate=None, max_bin=None,  
              max_cat_threshold=None, max_cat_to_onehot=None,  
              max_delta_step=None, max_depth=None, max_leaves=None,  
              min_child_weight=None, missing=nan, monotone_constraints=None,  
              multi_strategy=None, n_estimators=None, n_jobs=None,  
              num_parallel_tree=None, random_state=None, ...)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

## Predicting the Test set results

```
In [58]: y_pred = classifier.predict(X_test)  
y_pred  
  
array([0, 1, 0, ..., 0, 0, 0])
```

```
In [59]: ## Making the Confusion Matrix
```

```
In [60]: from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred)  
print(cm)  
  
[[1595  0]  
 [ 0  405]]
```

```
In [61]: from sklearn.metrics import accuracy_score  
ac = accuracy_score(y_test, y_pred)  
print(ac)
```

1.0

```
In [62]: bias = classifier.score(X_train,y_train)  
bias
```

1.0

```
In [63]: variance = classifier.score(X_test,y_test)  
variance
```

1.0

## Applying k-Fold Cross Validation

```
In [64]: from sklearn.model_selection import cross_val_score  
accuracies = cross_val_score(estimator = classifier, X = X_train, y = y_train, c  
print("Accuracy: {:.2f} %".format(accuracies.mean()*100))  
print("Standard Deviation: {:.2f} %".format(accuracies.std()*100))
```

Accuracy: 100.00 %  
Standard Deviation: 0.00 %

```
In [65]: # Training the K-NN model on the Training set  
from sklearn.neighbors import KNeighborsClassifier  
classifier = KNeighborsClassifier()  
classifier.fit(X_train, y_train)
```

KNeighborsClassifier()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [66]: # Predicting the Test set results  
y_pred = classifier.predict(X_test)  
y_pred  
  
array([0, 1, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [67]: # Making the Confusion Matrix  
from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred)  
print(cm)  
  
[[1595  0]  
 [ 0 405]]
```

```
In [68]: from sklearn.metrics import accuracy_score  
ac = accuracy_score(y_test, y_pred)  
print(ac)  
  
1.0
```

```
In [69]: bias = classifier.score(X_train,y_train)  
bias  
  
1.0
```

```
In [70]: variance = classifier.score(X_test,y_test)  
variance  
  
1.0
```

```
In [71]: # train a Gaussian Naive Bayes classifier on the training set
from sklearn.naive_bayes import GaussianNB

# instantiate the model
gnb = GaussianNB()

# fit the model
gnb.fit(X_train, y_train)
```

```
GaussianNB()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [72]: y_pred = gnb.predict(X_test)
```

```
y_pred
```

```
array([0, 1, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [73]: from sklearn.metrics import accuracy_score
```

```
print('Model accuracy score: {:.4f}'.format(accuracy_score(y_test, y_pred)))
```

```
Model accuracy score: 1.0000
```

```
In [74]: # print the scores on training and test set
```

```
print('Training set score: {:.4f}'.format(gnb.score(X_train, y_train)))
```

```
print('Test set score: {:.4f}'.format(gnb.score(X_test, y_test)))
```

```
Training set score: 1.0000
```

```
Test set score: 1.0000
```

```
In [75]: bias = classifier.score(X_train,y_train)  
bias  
  
1.0
```

```
In [76]: variance = classifier.score(X_test,y_test)  
variance  
  
1.0
```

```
In [77]: # Training the SVM model on the Training set  
from sklearn.svm import SVC  
classifier = SVC()  
classifier.fit(X_train, y_train)  
  
SVC()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [78]: # Predicting the Test set results  
y_pred = classifier.predict(X_test)  
y_pred  
  
array([0, 1, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [79]: # Making the Confusion Matrix  
from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred)  
print(cm)  
  
[[1595  0]  
 [ 0  405]]
```

```
In [80]: # This is to get the Models Accuracy  
from sklearn.metrics import accuracy_score  
ac = accuracy_score(y_test, y_pred)  
print(ac)
```

1.0

```
In [81]: bias = classifier.score(X_train,y_train)  
bias
```

1.0

```
In [82]: variance = classifier.score(X_test,y_test)  
variance
```

1.0

```
In [83]: # This is to get the Classification Report  
from sklearn.metrics import classification_report  
cr = classification_report(y_test, y_pred)  
cr
```

	precision	recall	f1-score	support	0	1.00	1.00	1.00
accuracy	1.00	405\n	accuracy		1.00	2000\n	macro avg	
weighted avg	1.00	1.00	1.00	2000\n				

```
In [84]: # Training the Logistic Regression model on the Training set
```

```
from sklearn.linear_model import LogisticRegression  
classifier = LogisticRegression()  
classifier.fit(X_train, y_train)
```

```
LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [85]: # Predicting the Test set results  
y_pred = classifier.predict(X_test)  
y_pred  
  
array([0, 1, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [86]: # Making the Confusion Matrix  
from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred)  
print(cm)  
  
[[1595 0]  
 [ 0 405]]
```

```
In [87]: # This is to get the Models Accuracy  
from sklearn.metrics import accuracy_score  
ac = accuracy_score(y_test, y_pred)  
print(ac)  
  
1.0
```

```
In [88]: # This is to get the Classification Report  
from sklearn.metrics import classification_report  
cr = classification_report(y_test, y_pred)  
cr  
  
'precision    recall    f1-score    support\n\n          1.00      405\\n\\n      accuracy           1.00      2000\\n\\n  
weighted avg      1.00      1.00      1.00      2000\\n'
```

```
In [89]: bias = classifier.score(X_train, y_train)  
bias  
  
1.0
```

```
In [90]: variance = classifier.score(X_test, y_test)  
variance  
  
1.0
```

```
In [91]: # print the scores on training and test set

print('Training set score: {:.4f}'.format(gnb.score(X_train, y_train)))

print('Test set score: {:.4f}'.format(gnb.score(X_test, y_test)))

Training set score: 1.0000
Test set score: 1.0000
```

```
In [ ]:
```